

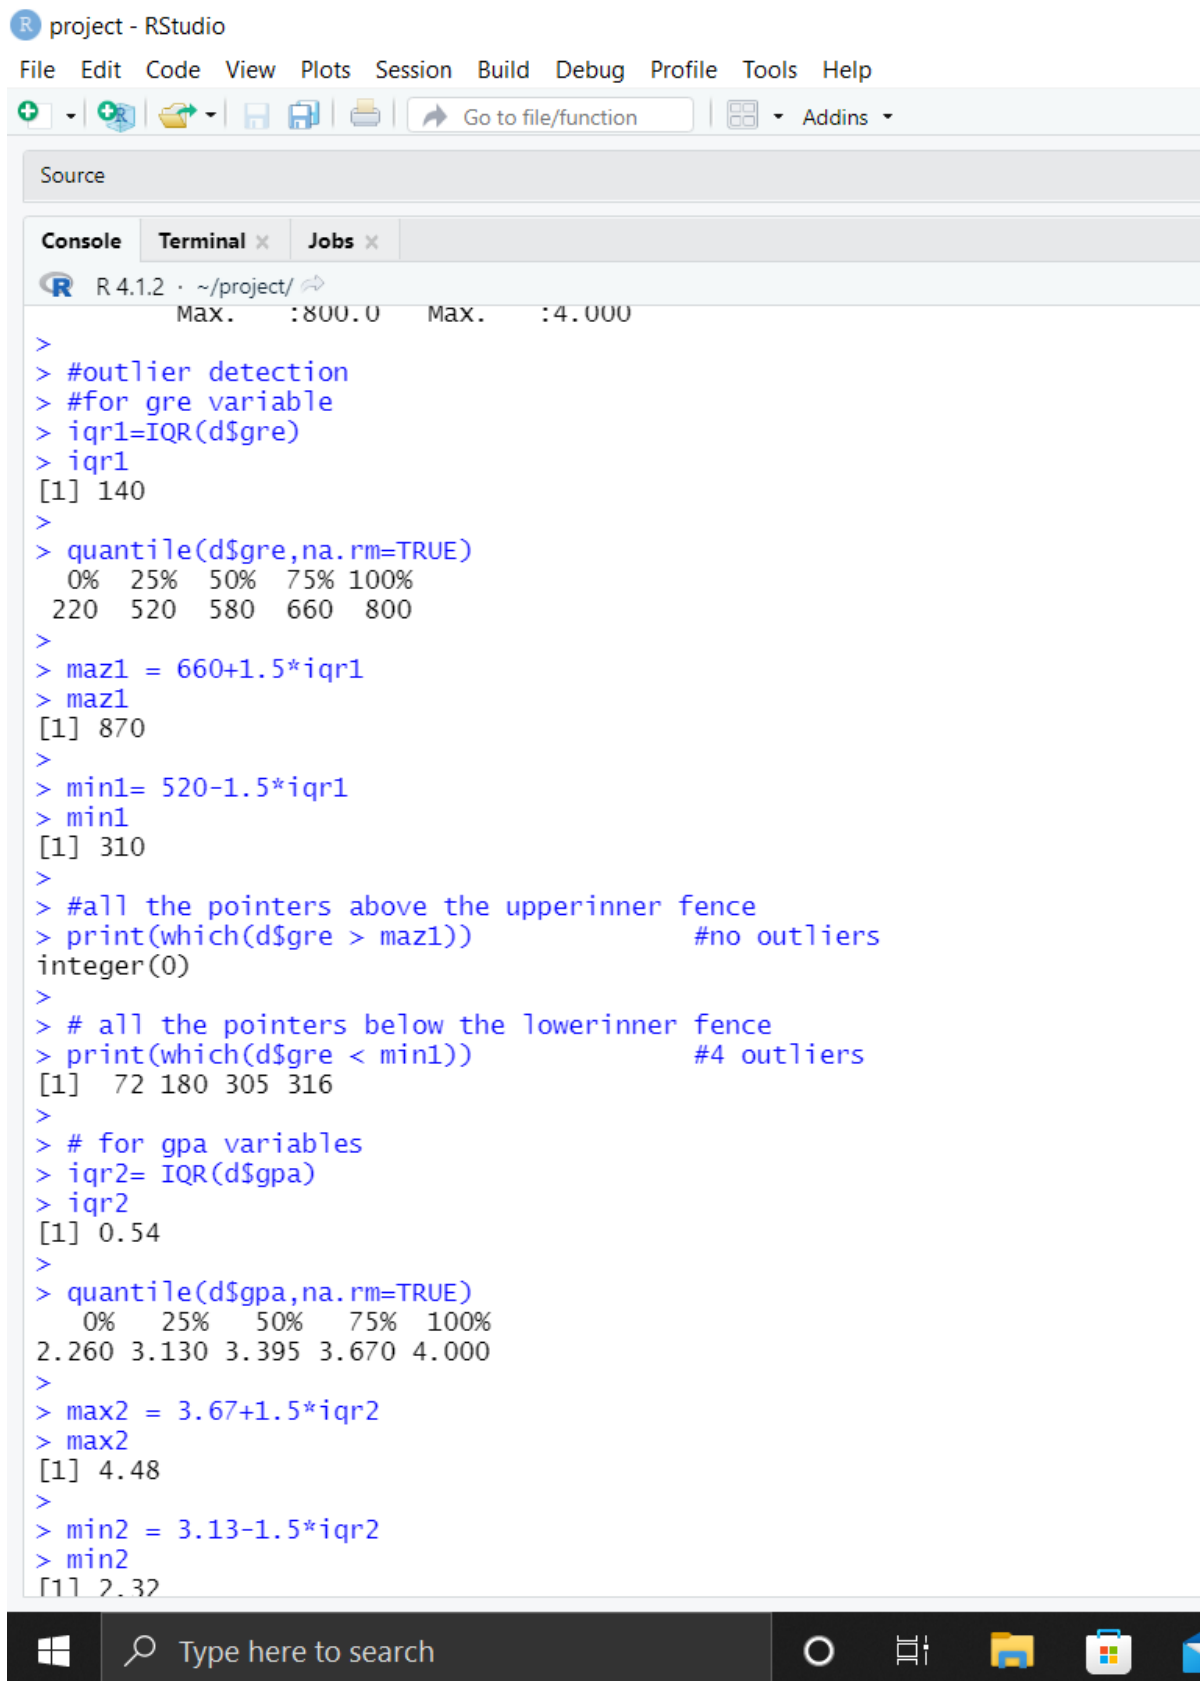
Find the missing values. (if any, perform missing value treatment)

```
project - RStudio
File Edit Code View Plots Session Build Debug Profile Tools Help
Go to file/function Addins

Source

Console Terminal Jobs
R 4.1.2 · ~/project/
> #importing college admission data set
>
> d = read.csv("C:/Users/ADMIN/Documents/project/College_admission.csv",header=T)
> head(d)
  admit gre  gpa ses Gender_Male Race rank
1     0 380 3.61  1           0     3     3
2     1 660 3.67  2           0     2     3
3     1 800 4.00  2           0     2     1
4     1 640 3.19  1           1     2     4
5     0 520 2.93  3           1     2     4
6     1 760 3.00  2           1     1     2
> d1 = d
>
> #no of records
> nrow(d)
[1] 400
>
> #datatype of columns
> sapply(d, class)
      admit      gre      gpa      ses Gender_Male      Race      rank
"integer" "integer" "numeric" "integer" "integer" "integer" "integer"
>
> #converting columns to factor type
> cols = c(1,4:7)
> d[,cols]<-lapply(d[,cols], factor)
> sapply(d,class)
      admit      gre      gpa      ses Gender_Male      Race      rank
"factor"  "integer" "numeric" "factor"  "factor"  "factor"  "factor"
>
> #count of missing values
> sum(is.na(d))
[1] 0
>
> #summary of dataset
> summary(d)
admit      gre      gpa      ses  Gender_Male  Race      rank
0:273   Min.   :220.0   Min.   :2.260   1:132   0:210      1:143   1: 61
1:127   1st Qu.:520.0   1st Qu.:3.130   2:139   1:190      2:129   2:151
        Median :580.0   Median :3.395   3:129           3:128   3:121
        Mean   :587.7   Mean    :3.390           4: 67
        3rd Qu.:660.0   3rd Qu.:3.670
        Max.   :800.0   Max.    :4.000
>
> #outlier detection
```

Find outliers (if any, then perform outlier treatment)



The screenshot shows the RStudio interface with the following components:

- Top Bar:** File Edit Code View Plots Session Build Debug Profile Tools Help
- Source Panel:** Empty
- Console Panel:** Contains the following R code and output:

```
>
> #outlier detection
> #for gre variable
> iqr1=IQR(d$gre)
> iqr1
[1] 140
>
> quantile(d$gre,na.rm=TRUE)
 0%  25%  50%  75% 100%
220  520  580  660  800
>
> maz1 = 660+1.5*iqr1
> maz1
[1] 870
>
> min1= 520-1.5*iqr1
> min1
[1] 310
>
> #all the pointers above the upperinner fence
> print(which(d$gre > maz1))           #no outliers
integer(0)
>
> # all the pointers below the lowerinner fence
> print(which(d$gre < min1))           #4 outliers
[1]  72 180 305 316
>
> # for gpa variables
> iqr2= IQR(d$gpa)
> iqr2
[1] 0.54
>
> quantile(d$gpa,na.rm=TRUE)
 0%  25%  50%  75% 100%
2.260 3.130 3.395 3.670 4.000
>
> max2 = 3.67+1.5*iqr2
> max2
[1] 4.48
>
> min2 = 3.13-1.5*iqr2
> min2
[1] 2.32
```
- Bottom Bar:** Windows taskbar with search bar and icons for RStudio, File Explorer, and other applications.

Console Terminal x Jobs x

R 4.1.2 · ~/project/ ↗

```
>
> # all the pointers above the upperinner fence
> print(which(d$gpa > max2))          # no outlier
integer(0)
>
> print(which(d$gpa < min2))          # 1 outlier
[1] 290
>
> #removing outliers
>
> d=d[-c(72, 180, 290, 305, 316),]
> nrow(d)
[1] 395
.
```

Splitting the data into train and test

```
> #splitting of data set into train and test
>
> set.seed(0)
> library("caTools")
>
> set.seed(0)
> library("caTools")
> d[,2:3]=scale(d[,2:3])
> split=sample.split(d$admit,SplitRatio = .75)
> train=subset(d,split==T)
> test=subset(d,split==F)
>
```

Run logistic model to determine the factors that influence the admission process of a student (Drop insignificant variables)

```
project - RStudio
File Edit Code View Plots Session Build Debug Profile Tools Help
Go to file/function Addins

Source

Console Terminal Jobs
R 4.1.2 · ~/project/

> #logistic regression
> logit1=glm(admit~.,train,family='binomial') #all variables
> summary(logit1)

Call:
glm(formula = admit ~ ., family = "binomial", data = train)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.8458  -0.8294  -0.5794   0.9459   2.1850

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept)  1.03714    0.45482   2.280  0.02259 *
gre          0.27452    0.15262   1.799  0.07206 .
gpa          0.51793    0.16125   3.212  0.00132 **
ses2        -0.38459    0.33468  -1.149  0.25050
ses3        -0.40768    0.34356  -1.187  0.23538
Gender_Male1 -0.09887    0.27541  -0.359  0.71959
Race2        -0.34389    0.33981  -1.012  0.31153
Race3        -0.43592    0.33303  -1.309  0.19054
rank2        -1.29613    0.40854  -3.173  0.00151 **
rank3        -1.70399    0.43413  -3.925  8.67e-05 ***
rank4        -2.05159    0.51218  -4.006  6.19e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 370.01  on 295  degrees of freedom
Residual deviance: 319.50  on 285  degrees of freedom
AIC: 341.5

Number of Fisher Scoring iterations: 4

>
> # in the above model gre,ses,gender_male and race variable are not significant
> # building new model with only gpa rank variables
>
> logit2=glm(admit~gpa+rank,train,family = 'binomial') # all the variables
> summary(logit2)

Call:
```

Second logistic model by removing insignificant variables

```
project - RStudio
File Edit Code View Plots Session Build Debug Profile Tools Help
+ [Go to file/function] [Addins]

Source

Console Terminal x Jobs x
R 4.1.2 · ~/project/

>
> # in the above model gre,ses,gender_male and race variable are not significant
> # building new model with only gpa rank variables
>
> logit2=glm(admit~gpa+rank,train,family = 'binomial')      # all the variables
> summary(logit2)

Call:
glm(formula = admit ~ gpa + rank, family = "binomial", data = train)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.8203  -0.8527  -0.5997   1.0019   2.2480

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept)   0.4765     0.3386   1.407  0.15931
gpa            0.6037     0.1479   4.081 4.49e-05 ***
rank2         -1.2261     0.3978  -3.082  0.00205 **
rank3         -1.7363     0.4248  -4.087 4.37e-05 ***
rank4         -2.0552     0.5038  -4.079 4.52e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 370.01  on 295  degrees of freedom
Residual deviance: 326.90  on 291  degrees of freedom
AIC: 336.9

Number of Fisher Scoring iterations: 4

>
> # Accuracy of logit model
>
> predicted_val1 = predict(logit1,test,type="response")
>
> test$pred_admit1= ifelse(predicted_val1>0.5,1,0)
>
> #confusion matrix
>
> conf_mat1=table(predicted=test$pred_admit1,actual=test$admit)
> conf_mat1
```

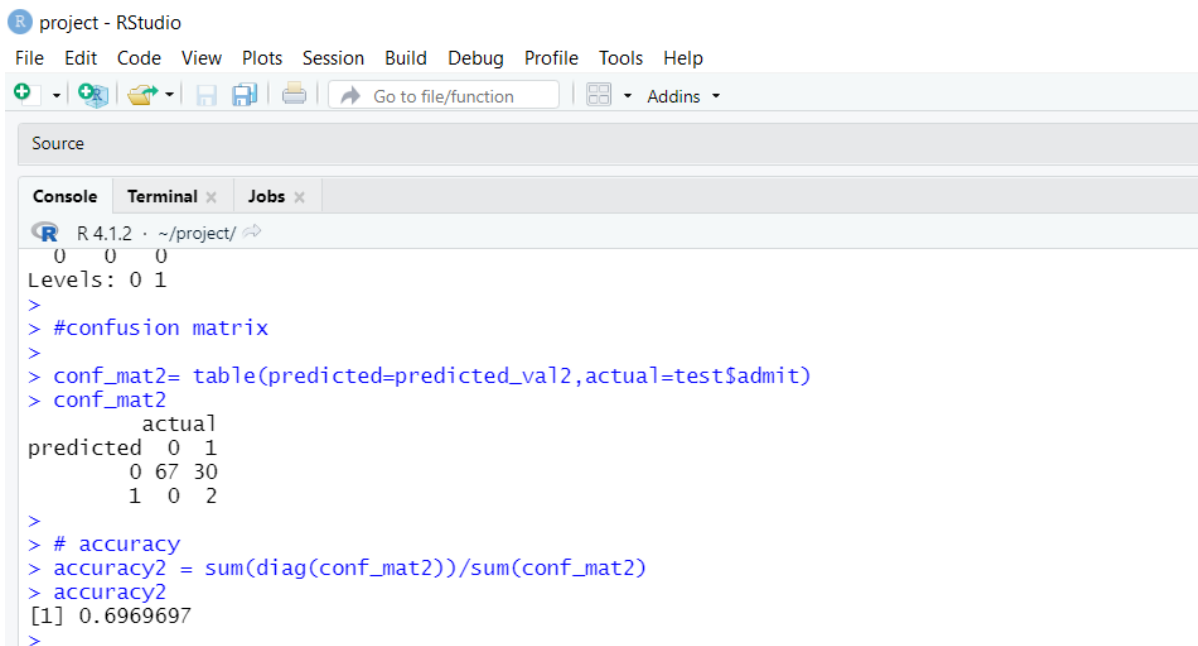
Here residual deviation increases so we will use first model

Accuracy of Logistic model

```
>
> # Accuracy of logit model
>
> predicted_val1 = predict(logit1,test,type="response")
>
> test$pred_admit1= ifelse(predicted_val1>0.5,1,0)
>
> #confusion matrix
>
> conf_mat1=table(predicted=test$pred_admit1,actual=test$admit)
> conf_mat1
      actual
predicted 0  1
      0 55 26
      1 12  6
>
> # accuracy
> accuracy1 = sum(diag(conf_mat1))/sum(conf_mat1)
> accuracy1
[1] 0.6161616
>
```

```
project - RStudio  
File Edit Code View Plots Session Build Debug Profile Tools Help  
 Go to file/function Addins  
Source  
Console Terminal x Jobs x  
 R 4.1.2 · ~/project/  
> #svm  
>  
> library(e1071)  
> svm_clf = svm(admit ~ . ,train,type = 'C-classification', kernel = 'linear')  
> summary(svm_clf)  
  
Call:  
svm(formula = admit ~ ., data = train, type = "C-classification", kernel = "linear")  
  
Parameters:  
SVM-Type: C-classification  
SVM-Kernel: radial  
cost: 1  
  
Number of Support Vectors: 207  
( 94 113 )  
  
Number of Classes: 2  
  
Levels:  
0 1  
  
>  
> #accuracy of svm  
>  
> predicted_val2 = predict(svm_clf,test[-1])  
> predicted_val2  
   1  11  14  16  26  29  31  35  37  38  52  60  61  63  68  70  74  82  94  95 102 104 109 113  
    0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0  
114 116 118 121 126 127 137 138 139 144 150 151 159 161 172 176 179 192 197 198 202 203 205 206  
    0   0   0   0   0   0   0   0   0   0   0   1   0   0   0   0   0   0   0   0   0   1   0   0  
212 214 215 216 223 233 236 238 243 247 248 251 253 256 260 266 269 270 274 276 277 279 285 286  
    0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0  
288 294 296 304 312 315 317 320 321 324 325 332 339 342 358 359 369 370 373 375 382 385 386 391  
    0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0   0  
394 396 400  
    0   0   0  
Levels: 0 1
```


Accuracy of SVM Model



The screenshot shows the RStudio interface with the console pane active. The console displays the following R code and output:

```
R 4.1.2 · ~/project/ ↗  
0 0 0  
Levels: 0 1  
>  
> #confusion matrix  
>  
> conf_mat2= table(predicted=predicted_val2,actual=test$admit)  
> conf_mat2  
      actual  
predicted 0  1  
      0 67 30  
      1  0  2  
>  
> # accuracy  
> accuracy2 = sum(diag(conf_mat2))/sum(conf_mat2)  
> accuracy2  
[1] 0.6969697  
>
```

The confusion matrix shows 67 true positives, 30 false positives, 0 false negatives, and 2 true negatives. The calculated accuracy is approximately 0.697.

Decision tree model

```
project - RStudio
File Edit Code View Plots Session Build Debug Profile Tools Help
Go to file/function Addins

Source

Console Terminal x Jobs x
R 4.1.2 · ~/project/

>
> ## decision tree
>
> library("rpart")
> library("rpart.plot")
> nrow(train)
[1] 296
> nrow(test)
[1] 99
> 0.03*nrow(train)
[1] 8.88
> 0.03*nrow(train)*3
[1] 26.64
>
> r.centrl=rpart.control(minsplit = 26,minbucket=9 , xavl = 5)
> dec_clf = rpart(admit~.,control =r.centrl,data=train)
> rpart.plot(dec_clf)
> summary(dec_clf)
Call:
rpart(formula = admit ~ ., data = train, control = r.centrl)
n= 296

      CP nsplit rel error      xerror      xstd
1 0.11702128      0 1.0000000 1.0000000 0.08520516
2 0.05319149      1 0.8829787 0.9574468 0.08419400
3 0.02127660      2 0.8297872 0.9680851 0.08445472
4 0.01063830      6 0.7446809 0.9680851 0.08445472
5 0.01000000      8 0.7234043 1.0000000 0.08520516

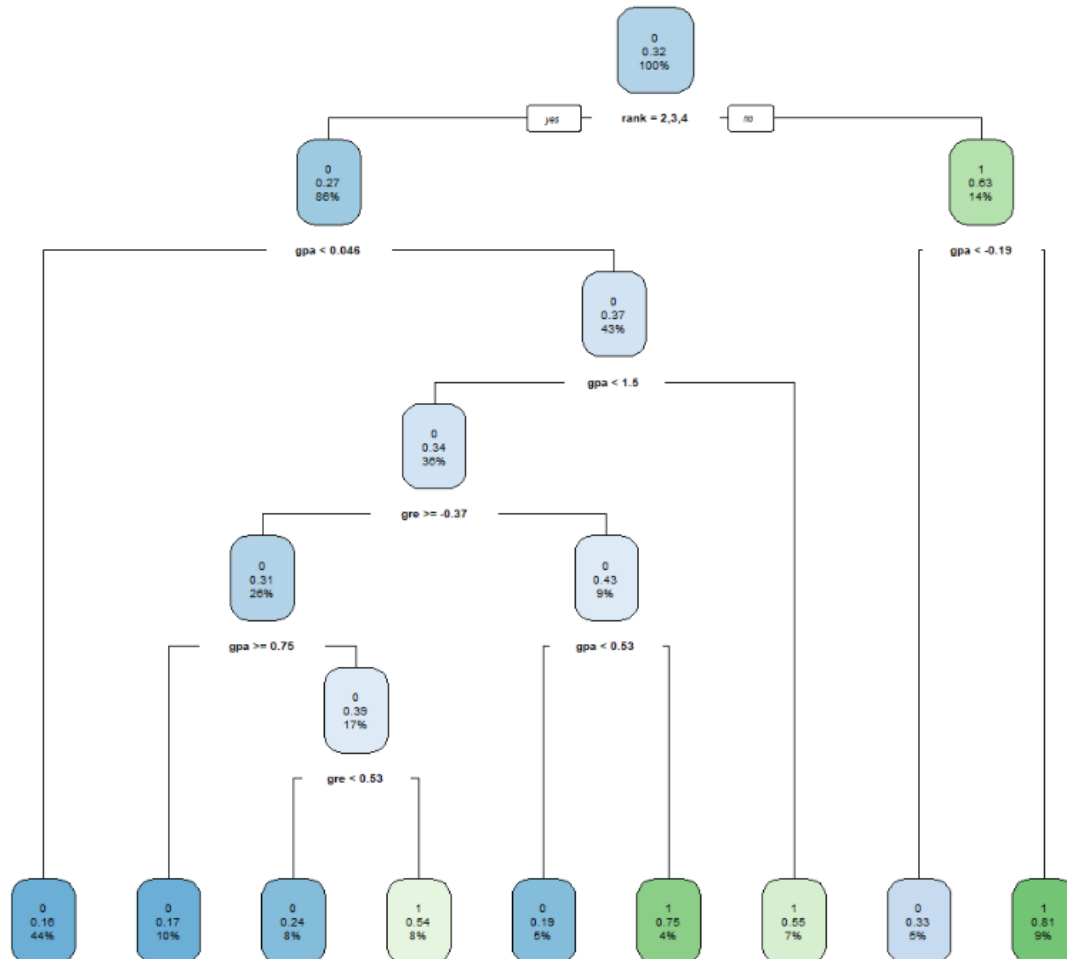
Variable importance
      gpa      rank      gre      Race Gender_Male      ses
      46       28       17       4         3         2

Node number 1: 296 observations,      complexity param=0.1170213
predicted class=0 expected loss=0.3175676 P(node) =1
class counts: 202 94
probabilities: 0.682 0.318
left son=2 (255 obs) right son=3 (41 obs)
Primary splits:
rank splits as RLLL, improve=9.5395740, (0 missing)
gpa < 0.1259629 to the left, improve=9.3487160, (0 missing)
gre < -0.7279174 to the left, improve=4.7657660, (0 missing)
ses splits as RLL, improve=1.0541450, (0 missing)
Race splits as RLL, improve=0.8225416, (0 missing)
```

Decision tree

Plot Zoom

— □ ×



Accuracy of Decision Tree

project - RStudio

File Edit Code View Plots Session Build Debug Profile Tools Help

Go to file/function Addins

Source

Console Terminal Jobs

R 4.1.2 · ~/project/

```

probabilities: 0.458 0.542

>
> #accuracy of decision tree
>
> predicted_val3 = predict(dec_clf, test[-1], type="class")
> predicted_val3
 1  11  14  16  26  29  31  35  37  38  52  60  61  63  68  70  74  82  94  95 102 104 109 113
114 116 118 121 126 127 137 138 139 144 150 151 159 161 172 176 179 192 197 198 202 203 205 206
0  1  0  0  1  0  1  0  1  0  0  0  1  1  0  0  0  0  1  0  0  1  1  0  0
212 214 215 216 223 233 236 238 243 247 248 251 253 256 260 266 269 270 274 276 277 279 285 286
0  0  0  0  0  0  0  1  0  0  0  0  1  0  0  0  1  0  0  0  1  0  0  0  0
288 294 296 304 312 315 317 320 321 324 325 332 339 342 358 359 369 370 373 375 382 385 386 391
0  1  0  1  1  0  0  0  0  0  0  1  0  0  0  0  1  0  0  0  0  0  0  0  0
394 396 400
0  1  0
Levels: 0 1
>
> #confusion matrix
>
> conf_mat3 = table(predicted=predicted_val3, actual=test$admit)
> conf_mat3
      actual
predicted 0  1
      0 50 22
      1 17 10
>
> #accuracy
>
> accuracy3= sum(diag(conf_mat3))/sum(conf_mat3)
> accuracy3
[1] 0.6060606

```

KNN and its Accuracy

```
project - RStudio
File Edit Code View Plots Session Build Debug Profile Tools Help
+ [icons] Go to file/function [Addins]
Source
Console Terminal x Jobs x
R 4.1.2 · ~/project/
> accuracy3
[1] 0.6060606
>
> #KNN and its accuracy
>
> library("class")
> knn = knn(train, test[-1],train$admit, k=19)
> knn
[1] 0 1 1 0 1 1 0 0 0 0 0 0 1 0 0 1 0 1 1 1 0 0 0 0 0 1 1 0 0 0 0 1 1 1 1 0 1 1 1 1 0 0 1 0 0
[48] 1 1 1 1 1 0 0 1 0 1 1 0 1 0 0 0 0 1 0 1 0 0 1 0 0 1 1 1 1 0 0 0 0 0 0 0 0 0 1 0 0 1 0 0 1 0
[95] 0 1 0 0 0
Levels: 0 1
>
> #confusion matrix
>
> conf_mat4=table(predicted=knn,actual=test$admit)
> conf_mat4
      actual
predicted 0  1
      0 41 15
      1 26 17
>
> # accuracy
>
> accuracy4= sum(diag(conf_mat4))/sum(conf_mat4)
> accuracy4
[1] 0.5858586
```

Naïve Bayes

project - RStudio

File Edit Code View Plots Session Build Debug Profile Tools Help

Go to file/function Addins

Source

Console Terminal Jobs

```
R 4.1.2 · ~/project/
> accuracy
[1] 0.5858586
>
> #naive bayes
>
> nb=naiveBayes(admit~. ,data=train)
> nb
```

Naive Bayes Classifier for Discrete Predictors

Call:
naiveBayes.default(x = X, y = Y, laplace = laplace)

A-priori probabilities:

Y

	0	1
0	0.6824324	0.3175676

Conditional probabilities:

gre

Y	[,1]	[,2]
0	-0.1406362	1.0036519
1	0.2862818	0.9187608

gpa

Y	[,1]	[,2]
0	-0.1601170	0.9929025
1	0.3759513	0.8978631

ses

Y	1	2	3
0	0.2920792	0.3613861	0.3465347
1	0.3829787	0.3191489	0.2978723

Gender_Male

Y	0	1
0	0.5297030	0.4702970
1	0.5319149	0.4680851

Race

Y	1	2	3
0	0.3217822	0.3465347	0.3316832
1	0.4042553	0.2978723	0.2978723

Windows taskbar: Type here to search, File Explorer, Calendar, Mail (6), Chrome, PDF Reader, and other icons.

Accuracy of Naïve Bayes

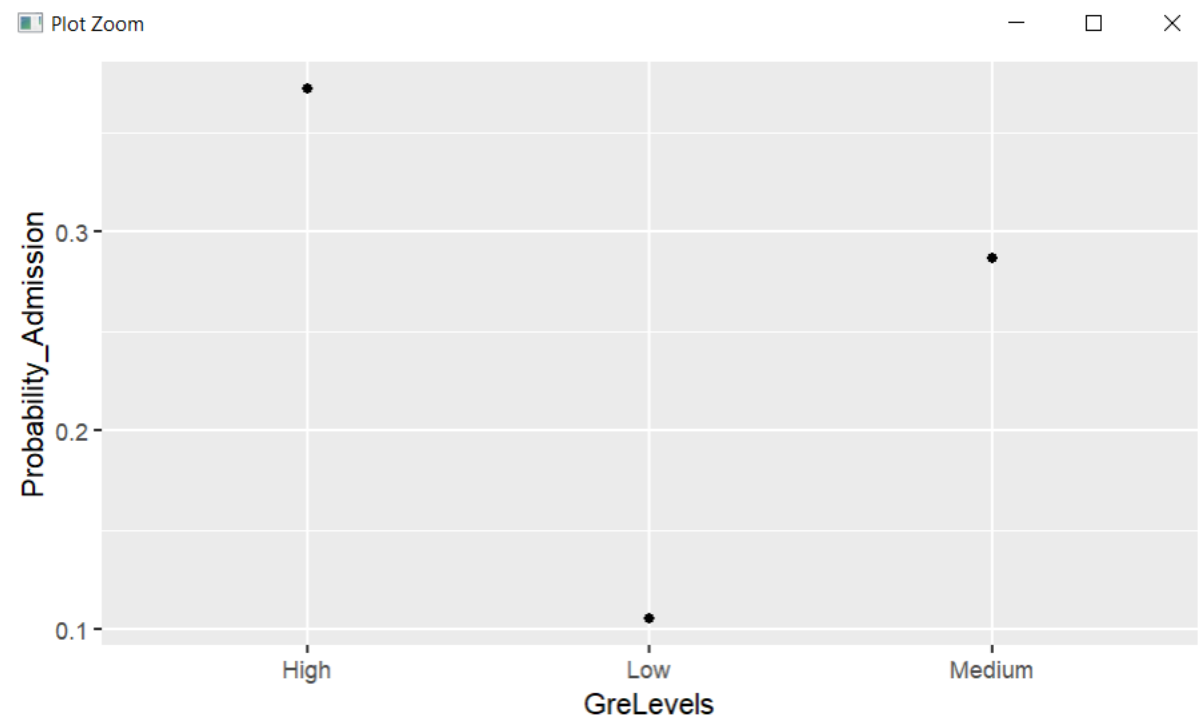
The screenshot shows the RStudio environment. At the top is the menu bar with options: File, Edit, Code, View, Plots, Session, Build, Debug, Profile, Tools, Help. Below the menu is a toolbar with icons for file operations and a search bar labeled "Go to file/function". The main window has tabs for "Source", "Console", "Terminal", and "Jobs". The "Source" tab is active, displaying an R script. The script contains several lines of code related to model evaluation, including calculating accuracy and creating a confusion matrix. The output of the script is visible in the console area at the bottom. The status bar at the very bottom indicates "R 4.1.2 · ~/project/" with a refresh icon.

```
>
> #accuracy of naive bayes
>
> predicted_val5 = predict(nb,test[-1],type="class")
> predicted_val5
[1] 0 1 0 0 1 1 0 0 1 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 1 0 0 0 0 1 0 1 0 0 1 1 0 0 0 0 0 0 0 0 0 0 1 1
[48] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 1 0 1 1 0 0 1 0 0 0 0 1 0 1 0 1 1 0 0 0 0
[95] 0 0 0 0 0
Levels: 0 1
>
> # confusion matrix
> conf_mat5=table(predicted=predicted_val5, actual=test$admit)
> conf_mat5
      actual
predicted 0 1
         0 52 25
         1 15 7
>
> # accuracy
> accuray=sum(diag(conf_mat5))/sum(conf_mat5)
> accuray
[1] 0.5959596
>
> ## logistic and svm are the best model with accuracy above 60%
> ### logistic model accuracy = 61.61%
> ## SVM model accuracy = 69.69%
```

Categorize the average of grade point into High, Medium, and Low (with admission probability percentages) and plot it on a point chart.

```
> #Descriptiven Categorize the average of grade point into High, Medium, and Low
>
> Descriptive = transform(d1,GreLevels=ifelse(gre<440,"Low",ifelse(gre<580,"Medium","High")))
> View(Descriptive)
> Sum_Desc=aggregate(admit~GreLevels,Descriptive,FUN=sum)
> length_Desc=aggregate(admit~GreLevels,Descriptive,FUN=length)
> Probability_table = cbind(Sum_Desc,Recs=length_Desc[,2])
> Probability_table_final = transform(Probability_table,Probability_Admission=admit/Recs)
> Probability_table_final
  GreLevels admit Recs Probability_Admission
1      High    84 226          0.3716814
2       Low     4  38          0.1052632
3    Medium    39 136          0.2867647
>
> library("ggplot2")
> ggplot(Probability_table_final,aes(x=GreLevels , y=Probability_Admission))+geom_point()
>
> #Cross grid for admission variable with GRE categorized
>
> table(Descriptive$admit,Descriptive$GreLevels)
      High Low Medium
0  142   34   97
1   84    4   39
>
>
>
> |
```

Point Chart



Descriptive showing GreLevels

project - RStudio

File Edit Code View Plots Session Build Debug Profile Tools Help

Go to file/function

Addins

Descriptive

Filter

	admit	gre	gpa	ses	Gender_Male	Race	rank	GreLevels
1	0	380	3.61	1	0	3	3	Low
2	1	660	3.67	2	0	2	3	High
3	1	800	4.00	2	0	2	1	High
4	1	640	3.19	1	1	2	4	High
5	0	520	2.93	3	1	2	4	Medium
6	1	760	3.00	2	1	1	2	High
7	1	560	2.98	2	1	2	1	Medium
8	0	400	3.08	2	0	2	2	Low
9	1	540	3.39	1	1	1	3	Medium
10	0	700	3.92	1	0	2	2	High
11	0	800	4.00	1	1	1	4	High
12	0	440	3.22	3	0	2	1	Medium
13	1	760	4.00	3	1	2	1	High
14	0	700	3.08	2	0	2	2	High
15	1	700	4.00	2	1	1	1	High
16	0	480	3.44	3	0	1	3	Medium
17	0	780	3.87	2	0	3	4	High
18	0	360	2.56	3	1	3	3	Low
19	0	800	3.75	1	1	3	2	High
20	1	540	3.81	1	0	3	1	Medium

Showing 1 to 20 of 400 entries, 8 total columns

Console

Windows Taskbar

Search: Type here to search

Taskbar Icons